Autonomy policy: first-response incident summaries

Day-one autonomy is L1 draft-only. During the first 30 minutes of an alert, the agent may read exported evidence, build a timeline, rank likely causes, and draft a human-verifiable update. It may not access live systems, restart services, alter config, roll back releases, or decide finance recovery.

L1 allowed: read alert/logs/metrics and cite file and line evidence for every claim; draft a short status note for human review and send.

Human-only actions: restarts, scaling, rule changes, rollbacks, replay, customer recovery, and any stakeholder message.

Grounding evidence: INC-2041 shows a release-to-symptom match and PayEx socket/timeouts after release 2.14.0. INC-2047 shows event-type drift from InvoiceApproved to invoice.approved while the worker subscription still filters on the old name.

Safety controls: exported evidence only; no live network tools or browser actions; output is DRAFT; uncited claims are discarded; a named human approves any move above L1.